

btctweets__sentiment__prediction-cnn-lstm-05312020

June 19, 2020

1 Bitcoin Tweets Sentiment Analysis Prediction – CNN-LSTM

```
[1]: # Data analysis and wrangling
import pandas as pd
import numpy as np
import os
import string
import csv

# Visualization
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn
from wordcloud import WordCloud

# Sentiment prediction
import re
import nltk
from sklearn.model_selection import train_test_split
from keras.preprocessing.text import Tokenizer
from keras.preprocessing import sequence

from tensorflow.keras import backend, models, layers
from tensorflow.keras.models import Model, Sequential
from tensorflow.keras.layers import Dense, Embedding, Conv1D, MaxPooling1D, LSTM
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report
```

Using TensorFlow backend.

```
[2]: # Upload the processed clean tweets

df = pd.read_csv('cleaned_tweet_data_05312020.csv',
                 header = 0,
                 error_bad_lines=False,
                 engine='python',
                 usecols=[2,5])
```

```
df.head()
```

```
[2]:          clean_text sentiment
0  goldman sachs hosting client call bitcoin gold...  neutral
1          ok president trump endorsing tommy  positive
2  bitcoin day nwhen btc first used commercial tr...  positive
3  great icle nif way could avoided know like min...  positive
4          bitcoin btc current price gbp  neutral
```

```
[3]: df.tail()
```

```
[3]:          clean_text sentiment
2015  xrp xrpcommunity btc bitcoin eth ltc vet oh ye...  positive
2016  yes world global pandemic world going recessio...  neutral
2017  last days grayscale bitcoin trust bought bitco...  neutral
2018  pretty incredible imho nwe need creativity lik...  positive
2019          drop sachs nget sats  neutral
```

1.1 Exploratory Data Analysis

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2020 entries, 0 to 2019
Data columns (total 2 columns):
clean_text    2020 non-null object
sentiment     2020 non-null object
dtypes: object(2)
memory usage: 31.7+ KB
```

```
[5]: df.describe()
```


```
[5]:          clean_text sentiment
count          2020          2020
unique          1391           3
top  tick tock less blocks  neutral
freq           36          964
```

```
[6]: tweets_count = df['sentiment'].value_counts()
tweets_count
```

```
[6]: neutral    964
positive    788
negative    268
Name: sentiment, dtype: int64
```

```
[7]: # Plot sentiment mood count
seaborn.countplot(x='sentiment', data=df)
plt.title('Bitcoin Sentiment Analysis - 2020')
plt.ylabel('Total Counts')
```

```
[7]: Text(0, 0.5, 'Total Counts')
```



```
[8]: # add text length column

df['text_length'] = df['clean_text'].apply(len)
```

```
[9]: df.head()
```

```
[9]:
```

| | clean_text | sentiment | text_length |
|---|---------------------------------------------------|-----------|-------------|
| 0 | goldman sachs hosting client call bitcoin gold... | neutral | 60 |
| 1 | ok president trump endorsing tommy | positive | 34 |
| 2 | bitcoin day nwhen btc first used commercial tr... | positive | 87 |
| 3 | great icle nif way could avoided know like min... | positive | 68 |
| 4 | bitcoin btc current price gbp | neutral | 29 |

```
[10]: df.tail()
```

```
[10]:
```

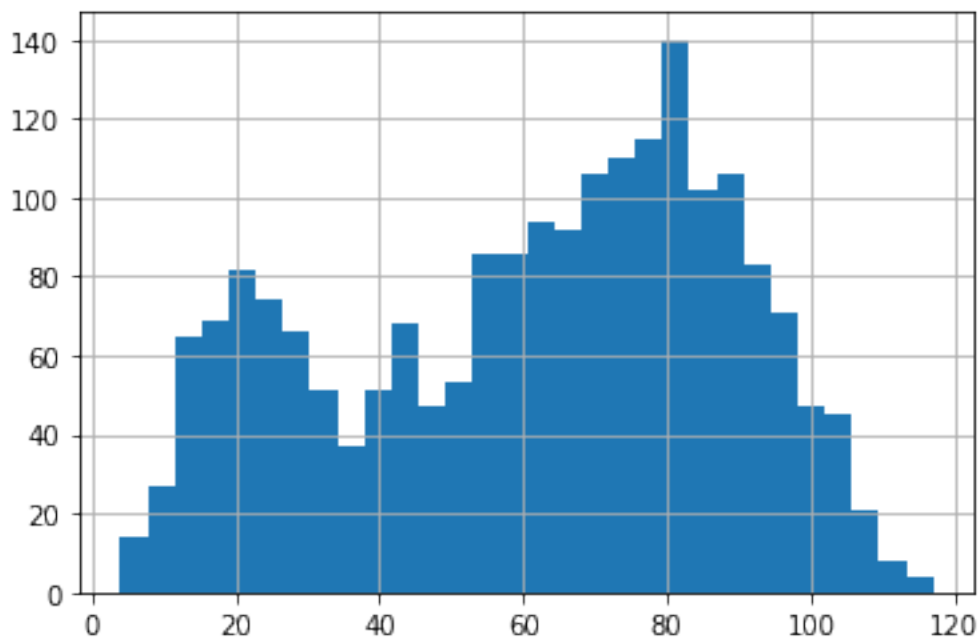
| | clean_text | sentiment | text_length |
|------|---------------------------------------------------|-----------|-------------|
| 2015 | xrp xrpcommunity btc bitcoin eth ltc vet oh ye... | positive | 96 |
| 2016 | yes world global pandemic world going recessio... | neutral | 56 |
| 2017 | last days grayscale bitcoin trust bought bitco... | neutral | 70 |
| 2018 | pretty incredible imho nwe need creativity lik... | positive | 68 |
| 2019 | drop sachs nget sats | neutral | 20 |

```
[11]: df['text_length'].describe()
```

```
[11]: count    2020.000000
      mean      61.235149
      std       27.054467
      min        4.000000
      25%       39.000000
      50%       66.500000
      75%       83.000000
      max      117.000000
      Name: text_length, dtype: float64
```

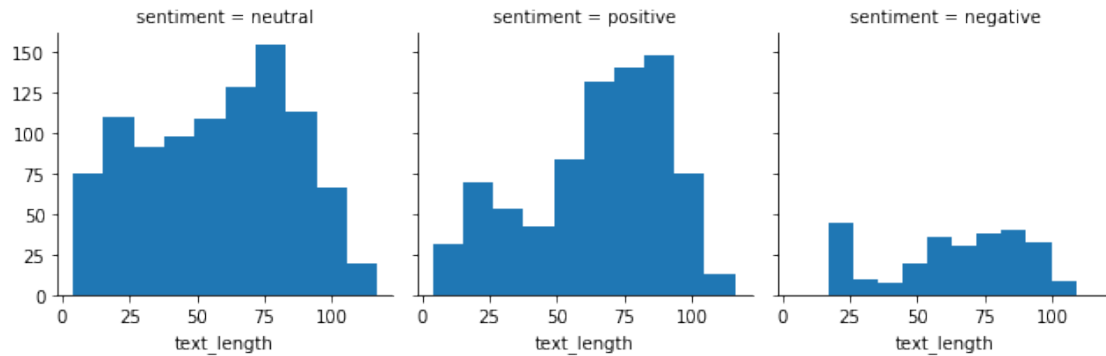
```
[12]: df['text_length'].hist(bins=30)
```

```
[12]: <matplotlib.axes._subplots.AxesSubplot at 0x219f61da548>
```



```
[13]: plot = seaborn.FacetGrid(df,col='sentiment')
      plot.map(plt.hist,'text_length')
```

```
[13]: <seaborn.axisgrid.FacetGrid at 0x219f7582648>
```



```
[14]: # visualization using wordcloud for the Neutral tweets
```

```
tweets_neutral = df[df['sentiment']=='neutral']
words = ' '.join(tweets_neutral['clean_text'])

wordcloud = WordCloud(background_color='blue').generate(words)

plt.figure(1,figsize=(15, 12))
plt.imshow(wordcloud)
plt.axis('off')
plt.show()
```



```
[15]: # visualization using wordcloud for the Positive tweets

tweets_positive = df[df['sentiment']=='positive']
words = ' '.join(tweets_positive['clean_text'])

wordcloud = WordCloud(background_color='orange').generate(words)

plt.figure(1,figsize=(15, 12))
plt.imshow(wordcloud)
plt.axis('off')
plt.show()
```




```
[16]: # visualization using wordcloud for the Negative tweets

tweets_negative = df[df['sentiment']=='negative']
words = ' '.join(tweets_negative['clean_text'])

wordcloud = WordCloud(background_color='green').generate(words)

plt.figure(1,figsize=(15, 12))
plt.imshow(wordcloud)
plt.axis('off')
plt.show()
```



1.2 Training and Testing sets

```
[17]: # Encode Categorical Variable for Deep Learning
      # Apply get_dummies

x = df['clean_text']
y = pd.get_dummies(df['sentiment']).values
num_classes = df['sentiment'].nunique()
```

```
[18]: # fix random seed for reproducibility
      np.random.seed(42)

      # Split into 80% train and 20% test sets
      x_train, x_test, y_train, y_test = train_test_split(x, y,
                                                          test_size=0.20,
                                                          stratify=y,
                                                          random_state=42)

      print(x_train.shape, x_test.shape, y_train.shape, y_test.shape)
```

(1616,) (404,) (1616, 3) (404, 3)

1.2.1 Tokenize Text

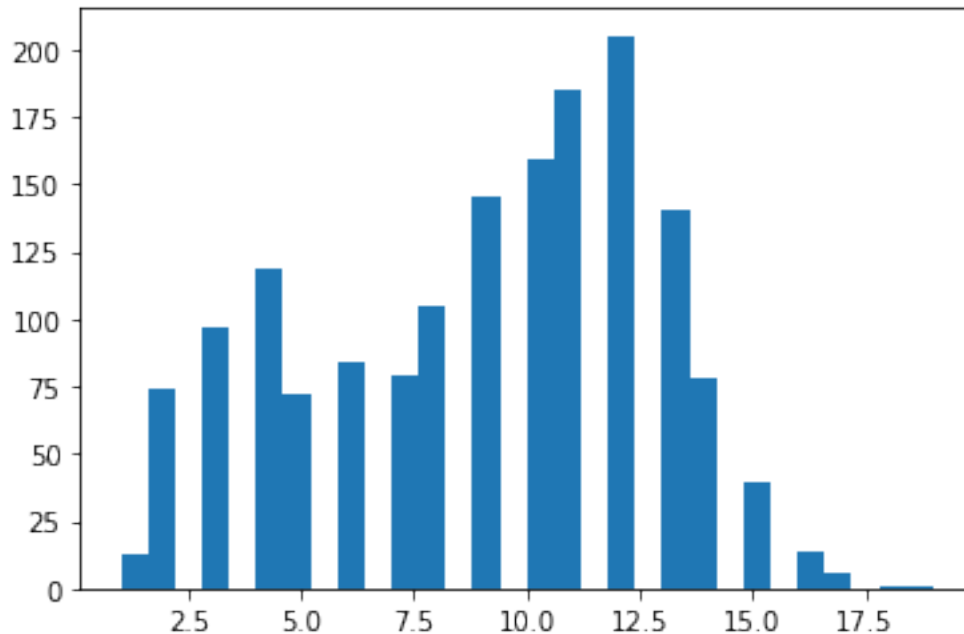
```
[19]: # Tokenize Text

      # Define max number of words in dictionary
      max_features = 20000
```

```
tokenizer = Tokenizer(num_words=max_features, lower=True)
tokenizer.fit_on_texts(list(x_train))

x_train = tokenizer.texts_to_sequences(x_train)
x_test = tokenizer.texts_to_sequences(x_test)
```

```
[20]: totalNumWords = [len(feature) for feature in x_train]
plt.hist(totalNumWords,bins = 30)
plt.show()
```



```
[21]: # Based on the plot histogram, setting feature maxWords
maxWords = 20

x_train = sequence.pad_sequences(x_train, maxlen=maxWords)
x_test = sequence.pad_sequences(x_test, maxlen=maxWords)

print(x_train.shape,x_test.shape)
```

```
(1616, 20) (404, 20)
```

1.3 Model – CNN Long Short-Term Memory Network (CNN LSTM)

```
[22]: # Create a CNN-LSTM Model

backend.clear_session()
```



```

embedding_size = 300
epoch = 20
batch_size = 128

model = models.Sequential()
model.add(Embedding(max_features, embedding_size, input_length=x_train.
    ↪shape[1]))
model.add(Conv1D(filters=32, kernel_size=3, padding='same', activation='relu'))
model.add(MaxPooling1D(pool_size=2))
model.add(Conv1D(filters=32, kernel_size=3, padding='same', activation='relu'))
model.add(MaxPooling1D(pool_size=2))
model.add(LSTM(96, dropout=0.15, recurrent_dropout=0.15))
model.add(Dense(num_classes, activation='softmax'))

# Compile the model
model.compile(optimizer = 'adam',
              loss = 'categorical_crossentropy',
              metrics = ['accuracy'])
print(model.summary())

# Train the model
history = model.fit(x_train,
                    y_train,
                    epochs = epoch,
                    batch_size = batch_size,
                    validation_data=(x_test, y_test),
                    verbose = 1)

history_dict = history.history
loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']
acc_values = history_dict['acc']
val_acc_values = history_dict['val_acc']
epochs = range(1, len(history_dict['acc']) + 1)

hist = pd.DataFrame(history.history)
print(hist.head())

plt.plot(epochs, loss_values, 'b', label = 'Training loss')
plt.plot(epochs, val_loss_values, 'r', label = 'Validation loss')
plt.title('Training and Validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()

```

```

plt.plot(epochs, acc_values, 'b', label = 'Training accuracy')
plt.plot(epochs, val_acc_values, 'r', label = 'Validation accuracy')
plt.title('Training and Validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

results = model.evaluate(x_test, y_test)
print(model.metrics_names)
print(results)

```

WARNING:tensorflow:From C:\cuong\lib\site-packages\tensorflow\python\ops\resource_variable_ops.py:435: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

WARNING:tensorflow:From C:\cuong\lib\site-packages\tensorflow\python\keras\backend.py:4010: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

| Layer (type) | Output Shape | Param # |
|--------------------------------|-----------------|---------|
| embedding (Embedding) | (None, 20, 300) | 6000000 |
| conv1d (Conv1D) | (None, 20, 32) | 28832 |
| max_pooling1d (MaxPooling1D) | (None, 10, 32) | 0 |
| conv1d_1 (Conv1D) | (None, 10, 32) | 3104 |
| max_pooling1d_1 (MaxPooling1D) | (None, 5, 32) | 0 |
| lstm (LSTM) | (None, 96) | 49536 |
| dense (Dense) | (None, 3) | 291 |

Total params: 6,081,763
 Trainable params: 6,081,763
 Non-trainable params: 0

None

Train on 1616 samples, validate on 404 samples

WARNING:tensorflow:From C:\cuong\lib\site-packages\tensorflow\python\ops\math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Epoch 1/20

1616/1616 [=====] - 6s 4ms/sample - loss: 1.0533 - acc: 0.4567 - val_loss: 0.9775 - val_acc: 0.4777

Epoch 2/20

1616/1616 [=====] - 3s 2ms/sample - loss: 0.9669 - acc: 0.5031 - val_loss: 0.9557 - val_acc: 0.5965

Epoch 3/20

1616/1616 [=====] - 4s 2ms/sample - loss: 0.8998 - acc: 0.6665 - val_loss: 0.8626 - val_acc: 0.6114

Epoch 4/20

1616/1616 [=====] - 3s 2ms/sample - loss: 0.6415 - acc: 0.7698 - val_loss: 0.6842 - val_acc: 0.6856

Epoch 5/20

1616/1616 [=====] - 3s 2ms/sample - loss: 0.3189 - acc: 0.8397 - val_loss: 0.7213 - val_acc: 0.7550

Epoch 6/20

1616/1616 [=====] - 4s 2ms/sample - loss: 0.1299 - acc: 0.9641 - val_loss: 0.9356 - val_acc: 0.7599

Epoch 7/20

1616/1616 [=====] - 3s 2ms/sample - loss: 0.0478 - acc: 0.9889 - val_loss: 0.9604 - val_acc: 0.7649

Epoch 8/20

1616/1616 [=====] - 3s 2ms/sample - loss: 0.0202 - acc: 0.9957 - val_loss: 1.2221 - val_acc: 0.7822

Epoch 9/20

1616/1616 [=====] - 3s 2ms/sample - loss: 0.0186 - acc: 0.9963 - val_loss: 0.9918 - val_acc: 0.7649

Epoch 10/20

1616/1616 [=====] - 3s 2ms/sample - loss: 0.0134 - acc: 0.9963 - val_loss: 1.1389 - val_acc: 0.7822

Epoch 11/20

1616/1616 [=====] - 4s 2ms/sample - loss: 0.0158 - acc: 0.9963 - val_loss: 1.0022 - val_acc: 0.7698

Epoch 12/20

1616/1616 [=====] - 4s 2ms/sample - loss: 0.0140 - acc: 0.9963 - val_loss: 1.0432 - val_acc: 0.7748

Epoch 13/20

1616/1616 [=====] - 3s 2ms/sample - loss: 0.0130 - acc: 0.9944 - val_loss: 1.0524 - val_acc: 0.7624

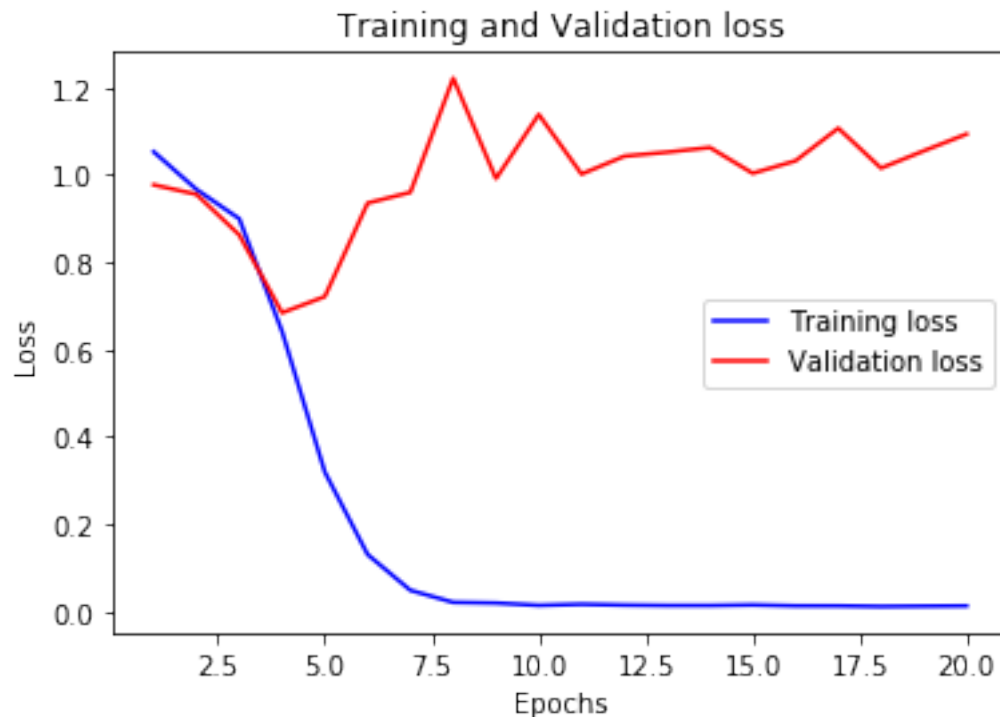
Epoch 14/20

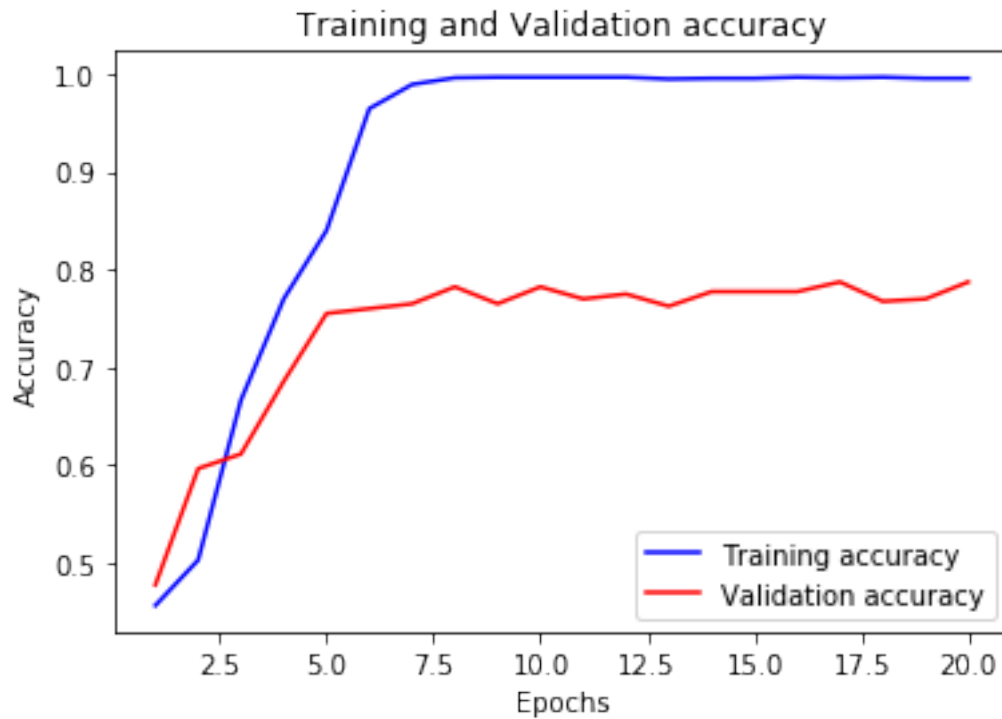
```

1616/1616 [=====] - 3s 2ms/sample - loss: 0.0130 - acc:
0.9950 - val_loss: 1.0631 - val_acc: 0.7772
Epoch 15/20
1616/1616 [=====] - 4s 2ms/sample - loss: 0.0143 - acc:
0.9950 - val_loss: 1.0039 - val_acc: 0.7772
Epoch 16/20
1616/1616 [=====] - 3s 2ms/sample - loss: 0.0122 - acc:
0.9963 - val_loss: 1.0325 - val_acc: 0.7772
Epoch 17/20
1616/1616 [=====] - 4s 2ms/sample - loss: 0.0121 - acc:
0.9957 - val_loss: 1.1074 - val_acc: 0.7871
Epoch 18/20
1616/1616 [=====] - 3s 2ms/sample - loss: 0.0107 - acc:
0.9963 - val_loss: 1.0158 - val_acc: 0.7673
Epoch 19/20
1616/1616 [=====] - 4s 2ms/sample - loss: 0.0114 - acc:
0.9950 - val_loss: 1.0554 - val_acc: 0.7698
Epoch 20/20
1616/1616 [=====] - 3s 2ms/sample - loss: 0.0121 - acc:
0.9950 - val_loss: 1.0940 - val_acc: 0.7871

```

| | loss | acc | val_loss | val_acc |
|---|----------|----------|----------|----------|
| 0 | 1.053291 | 0.456683 | 0.977466 | 0.477723 |
| 1 | 0.966943 | 0.503094 | 0.955698 | 0.596535 |
| 2 | 0.899789 | 0.666460 | 0.862560 | 0.611386 |
| 3 | 0.641451 | 0.769802 | 0.684168 | 0.685644 |
| 4 | 0.318946 | 0.839728 | 0.721333 | 0.754951 |





```
404/404 [=====] - 0s 274us/sample - loss: 1.0940 - acc:
0.7871
['loss', 'acc']
[1.0939605708169464, 0.7871287]
```

1.4 Model Evaluation

```
[23]: # Reference link: https://stackoverflow.com/questions/44189119/
      ↪ how-to-plot-confusion-matrix-correctly

      # Evaluate model with Test set

      def model_evaluation():
          # predict class with test set
          y_pred_test = model.predict_classes(x_test, batch_size=batch_size,
          ↪ verbose=0)
          print('Accuracy:\t{:.0.1f}%'.format(accuracy_score(np.
          ↪ argmax(y_test,axis=1),y_pred_test)*100))

          #classification report
          print('\n')
```

```

print(classification_report(np.argmax(y_test,axis=1), y_pred_test))

#confusion matrix
confmat = confusion_matrix(np.argmax(y_test,axis=1), y_pred_test)

fig, ax = plt.subplots(figsize=(4, 4))
ax.matshow(confmat, cmap=plt.cm.Blues, alpha=0.3)
for i in range(confmat.shape[0]):
    for j in range(confmat.shape[1]):
        ax.text(x=j, y=i, s=confmat[i, j], va='center', ha='center')
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.tight_layout()

```

```
[24]: print(model_evaluation())
```

Accuracy: 78.7%

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.88 | 0.53 | 0.66 | 53 |
| 1 | 0.76 | 0.88 | 0.81 | 193 |
| 2 | 0.81 | 0.77 | 0.79 | 158 |
| accuracy | | | 0.79 | 404 |
| macro avg | 0.81 | 0.72 | 0.75 | 404 |
| weighted avg | 0.79 | 0.79 | 0.78 | 404 |

None

